

Day 2 — Whiteboard Exercise Pack

Trainer reference for mid-hour activities · Half 1 + Half 2

HOW TO USE

Two activities — one per half. Both are 5-minute interludes designed to lift the room out of slide-watching mode. Keep this pack on the lectern.

Each exercise has a prompt to read aloud, an expected-answer set to collect, and a closeout to land the lesson.

Do not reveal answers before participants commit. The whole value of these drills is in the public defense.

CONTENTS

Exercise	When	Theme
1	After Slide 5 (~10:50)	Make this agentic
2	After Slide 19 (~12:30)	Framework decision drill

Make this agentic

Trigger participants to name the missing ingredients between a raw Bedrock call and a real agent. Pre-loads Slides 6–9.

PROMPT TO READ ALOUD

VERBATIM

"On the slide is the Bedrock call. Sketch on your notepad: what would you need to add to make this agentic? Name 3 things. We will compare in 2 minutes."

REFERENCE — THE BEDROCK CALL ON SCREEN

```
client.invoke_model(  
    modelId="anthropic.claude-3-5-sonnet-...",  
    body=json.dumps(  
        "system": "Confirm flight bookings.",  
        "messages": [{"role": "user", "content": "..."}]  
    )  
)
```

HOW TO RUN

00:00 – 00:30	Read the prompt. Make sure they understand the task — 'sketch on your notepad,' not 'shout out.'
00:30 – 02:30	Silent thinking time. 2 min. Walk the room — don't peek, just signal you're available.
02:30 – 04:30	Collect answers. Pick 4–5 hands. Write each on the whiteboard as it comes.
04:30 – 05:00	Closeout transition. Read the closeout line aloud, then click to Slide 6.

EXPECTED ANSWERS — COLLECT 4 OR 5

1. A **loop** — "call again with the result"
2. **Tool definitions** — "let it look up flight status, hotel availability, etc."
3. **Memory** — "remember the user across turns"
4. A **planner prompt** — "decide what to do next"

5. Stop conditions — "when do we stop looping"

Bonus answers (acknowledge but don't dwell): error handling, schema validation, observability, cost caps. Save these for Slide 8.

CLOSEOUT LINE

READ VERBATIM

"Hold those answers. Slide 6 is exactly that list — but on AWS. That is what AgentCore is."

Framework decision drill

Stress-test the rubric. Five scenarios. Each table picks Strands or LangGraph in 90 seconds. Trainer plays skeptic.

PROMPT TO READ ALOUD

VERBATIM

"Five scenarios coming up. Each table picks Strands or LangGraph. 90 seconds per scenario. Then one person at the table defends the choice in one sentence. Whichever framework you pick, I will push back — be ready to argue for the other one too. The rubric on Slide 19 is your friend."

THE FIVE SCENARIOS

1

Single-turn FAQ bot. 2 tools (knowledge-base lookup + escalate-to-human). No state needed across turns.

→ **Strands**

Why: no shared state, no HITL inside the loop, single-agent. Model-driven works.

2

Insurance claim triage. 6 steps end-to-end. Supervisor approval required before any payout over \$5,000.

→ **LangGraph**

Why: shared state across 6 steps + HITL gate. The decision tree stops at question 1.

3

Code review assistant. Reads a diff, comments inline on each file. Stateless between reviews.

→ **Strands**

Why: single-shot, stateless. Model loops over files and emits comments. No coordination needed.

4

Customer onboarding flow. 4 stages — verify identity, collect docs, run KYC, activate account. Can pause for days, resumes where left off.

→ **LangGraph**

Why: long-running checkpointing + shared state. Pausing for days mid-flow is the killer requirement.

5

Refund agent. Looks up order, computes refund amount, calls Stripe. Linear three-step.

→ **Strands**

Why: linear, no shared state, model can orchestrate the three calls confidently. Graph overhead unjustified.

HOW TO RUN

00:00 – 00:30	Read the prompt. Reference the rubric on Slide 19 — keep it on screen if possible.
00:30 – 04:00	Run all 5 scenarios. 30–40 seconds per scenario including pushback.
04:00 – 05:00	Score by show of hands. Don't dwell — the goal is internalizing the rubric, not winning.

SKEPTIC PUSHBACKS — TRAINER PLAYBOOK

Whichever framework a table picks, push back with the other. Sample pushbacks below — use them to force the defense.

If they pick Strands for scenario 2 (claim triage)

"How will you pause for the supervisor approval mid-loop? Show me what that looks like." → Forces them to realize the loop has to be cracked open.

If they pick LangGraph for scenario 1 (FAQ bot)

"You're building 7 nodes for a 2-tool bot. Why? What does the state schema buy you?" → Forces them to confront overengineering.

If they pick Strands for scenario 4 (onboarding)

"How does your agent resume when the user comes back 3 days later? Where does state live?" → Lands the checkpointing argument.

If they pick LangGraph for scenario 3 (code review)

"What state do you need shared? The diff goes in, comments come out. Where's the graph value?" → Lands that not everything needs a graph.

If they pick LangGraph for scenario 5 (refund)

"Three sequential calls. Do you really need typed state and edges? What if it grew to 5 steps?" → Probes the threshold thinking.

CLOSEOUT LINE

READ VERBATIM

"Three diagnostic questions, in order. Stop at the first YES. That is the rubric. Print it. Pin it. Refer to it on every new agent — yours and your team's."